

Assignment 2:

Memory task



LSTM and Recurrent Nets UE, 04.11.2019
Frederik Kratzert

Overview

- Assignment 1 Discussion
- Assignment 2

Assignment 1 Discussion

Gradient of bias terms

$$\mathbf{s}[t] = \mathbf{W}^T \mathbf{x}[t] + \mathbf{R}^T \mathbf{a}[t - 1] + \mathbf{b}_s$$

$$\mathbf{a}[t] = \tanh(\mathbf{s}[t])$$

$$\hat{\mathbf{y}} = \mathbf{V}^T \mathbf{a}[t = T] + \mathbf{b}_y$$

$$\frac{\partial L}{\partial \mathbf{b}_s[t]} = \dots$$

$$\frac{\partial L}{\partial \mathbf{b}_y[t]} = \dots$$

Weight initialization

```
def __init__(self, input_size: int, hidden_size: int, output_size: int):  
    """Initialization  
  
    Parameters  
    -----  
    input_size : int  
        Number of input features per time step  
    hidden_size : int  
        Number of hidden units in the RNN  
    output_size : int  
        Number of output units.  
    """  
    super(RNN, self).__init__()  
    self.input_size = input_size  
    self.hidden_size = hidden_size  
    self.output_size = output_size  
  
    # create and initialize weights of the network  
    self.W = np.random.uniform(-0.2, 0.2, (input_size, hidden_size))  
    self.R = np.random.uniform(-0.2, 0.2, (hidden_size, hidden_size))  
    self.bs = np.zeros((hidden_size, 1))  
    self.V = np.random.uniform(-0.2, 0.2, (hidden_size, output_size))  
    self.by = np.zeros((output_size, 1))
```

Forward pass

```
def forward(self, x: np.ndarray) -> np.ndarray:
    """Forward pass through the RNN.
```

```
    Parameters
```

```
    -----
```

```
    x : np.ndarray
```

```
        Input sequence(s) of shape [sequence length, number of features]
```

```
    Returns
```

```
    -----
```

```
    NumPy array containing the network prediction for the input sample.
```

```
    """
```

```
    seq_length, _ = x.shape
```

```
    self.a = np.zeros((seq_length+1, self.hidden_size))
```

← Store activations for backward pass

```
    a_0 = np.zeros((self.hidden_size))
```

```
    self.a[-1] = a_0.copy()
```

```
    for t in range(seq_length):
```

```
        s_t = np.dot(x[t], self.W) + np.dot(self.a[t-1], self.R) + self.bs.T
```

```
        a_t = np.tanh(s_t)
```

```
        self.a[t] = a_t
```

```
    y_hat = np.dot(a_t, self.V) + self.by.T
```

```
    self.y_hat = y_hat.copy()
```

```
    self.x = x.copy()
```

```
    return y_hat
```

$$s[t] = W^T x[t] + R^T a[t - 1] + b_s$$

$$a[t] = \tanh(s[t])$$

$$\hat{y} = V^T a[t = T] + b_y$$

Backward pass

```
def backward(self, d_loss: np.ndarray) -> Dict:
    """Calculate the backward pass through the RNN.

    Parameters
    -----
    d_loss : np.ndarray
        The gradient of the loss w.r.t the network output in the shape [output_size,]

    Returns
    -----
    Dictionary containing the gradients for each network weight as key-value pair.
    """
    seq_length, _ = self.x.shape

    # track gradients of weight matrices
    d_V = np.zeros_like(self.V)
    d_by = np.zeros_like(self.by)
    d_R = np.zeros((seq_length, *self.R.shape))
    d_W = np.zeros((seq_length, *self.W.shape))
    d_bs = np.zeros((seq_length, *self.bs.shape))
```

Backward pass

```
# track gradients of weight matrices
d_V = np.zeros_like(self.V)
d_by = np.zeros_like(self.by)
d_R = np.zeros((seq_length, *self.R.shape))
d_W = np.zeros((seq_length, *self.W.shape))
d_bs = np.zeros((seq_length, *self.bs.shape))

for t in reversed(range(seq_length)):
    if t == seq_length - 1:

        d_V = np.dot(self.a[t].reshape(-1,1), d_loss.reshape(1, -1))
        d_a = np.dot(d_loss.reshape(1, -1), self.V.T)
        d_by = d_loss.reshape(self.by.shape)
    else:
        d_a = d_a_next

    d_s = d_a * (1 - self.a[t]*self.a[t])
    d_W[t] = np.dot(self.x[t].reshape(-1,1), d_s)
    d_R[t] = np.dot(self.a[t-1].reshape(-1,1), d_s)
    d_bs[t] = d_s.reshape(d_bs[t].shape)

    d_a_next = np.dot(d_s, self.R.T)

self.grads = {'d_V': d_V, 'd_W': d_W, 'd_R': d_R, 'd_bs': d_bs, 'd_by': d_by}

return self.grads
```

$$\delta(t)^\top = \frac{\partial L}{\partial s(t)}$$

Parameter update

```
def update(self, lr: float):
    """Update the network parameter.

    Parameters
    -----
    lr : float
        Learning rate used for the weight update
    """
    if not self.grads:
        raise RuntimeError("You have to call the .backward() function first")

    for key, grad in self.grads.items():
        if len(grad.shape) == 3:
            self.grads[key] = grad.sum(axis=0)

    self.W -= lr*self.grads['d_W']
    self.R -= lr*self.grads['d_R']
    self.V -= lr*self.grads['d_V']
    self.bs -= lr*self.grads['d_bs']
    self.by -= lr*self.grads['d_by']

    # reset internal class attributes
    self.grads = {}
    self.y_hat, self.a = None, None
```

Numerical Gradient

$$\frac{\partial f}{\partial w_i} \approx \frac{f(x, w_1, \dots, w_i + \epsilon, \dots, w_n) - f(x, w_1, \dots, w_i - \epsilon, \dots, w_n)}{2 * \epsilon}$$

Numerical Gradient

```
weights = model.get_weights()
numerical_gradients = {key: np.zeros_like(val) for key, val in weights.items()}

for name, weight in weights.items():
    new_weights = {key: val for key, val in weights.items() if key != name}
    for i in range(weight.size):
        new_weight = weight.copy()
        m, n = np.unravel_index(i, weight.shape)

        # upper approximation
        new_weight = weight.copy()
        new_weight[m,n] += eps
        new_weights[name] = new_weight
        model.set_weights(new_weights)
        y_hat_upper = model.forward(x)

        # lower approximation
        new_weight = weight.copy()
        new_weight[m,n] -= eps
        new_weights[name] = new_weight
        model.set_weights(new_weights)
        y_hat_lower = model.forward(x)

        # calculate gradient approximation
        numerical_gradients[name][m,n] = np.sum((y_hat_upper - y_hat_lower) / (2*eps))

numerical_gradients = {f"d_{key}": val for key, val in numerical_gradients.items()}

return numerical_gradients
```

Analytical Gradient

```
def get_analytical_gradient(model: RNN, x: np.ndarray) -> Dict:
    """Helper function to get the analytical gradient.
```

Note: In contrast to the RNN update function, use the sum of the recurrent gradients over time (and average over batch samples)

Parameters

model : RNN

The RNN model object

x : np.ndarray

Input sequence(s) of shape [sequence length, number of features]

Returns

A dictionary containing the analytical gradients for each weight of the RNN.

"""

```
_ = model.forward(x)
```

```
analytical_grads = model.backward(np.ones((model.output_size)))
```

```
for key, grad in analytical_grads.items():
```

```
    if len(grad.shape) == 3:
```

```
        analytical_grads[key] = grad.sum(axis=0)
```

```
return analytical_grads
```

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial \hat{u}} \frac{\partial \hat{y}}{\partial w}$$

→ $\frac{\partial L}{\partial \hat{y}} = 1$

Assignment 2

Memory task



Tasks

- Implement data generator
 - Implement MSE loss (forward + backward)
 - Implement learner class
 - Train RNN for different sequence lengths
 - Visualize vanishing gradients
- RNN implementation is provided

Data Generator

```
def generate_samples(batch_size: int, seq_length: int) -> Tuple[np.ndarray, np.ndarray]:  
    """Data generator for memory task
```

Note: Implement this function as a Python generator.

Parameters

batch_size : int

Number of samples in one batch

seq_length : int

Length of sequence of random numbers

Returns

x : np.ndarray

Array of shape [sequence length, batch size, 1], where each sample is a sequence of random generated numbers between -1 and 1.

y : np.ndarray

Array of shape [batch size, 1], where each element i contains the label corresponding to sample i of the input array. The label is the first element of the sequence.

"""

```
#####  
# Your code comes here #  
#####
```


MSE Loss

- Forward pass:
$$L = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$
- Backward pass: $\longrightarrow \frac{\partial L}{\partial \hat{y}}$

Check (backward pass) implementation using numerical gradient check, again.

—————► Hint: This is a parameter free function, so you apply the variations directly to the model output.

Learner Class

- Used to facilitate training
 - Tracks loss
 - Trains model
 - Makes prediction
 - Plots results
- Use provided code skeleton and add missing code

Model Training

- What is the maximum sequence length the RNN can solve.
- For each length, train multiple repetitions.
- Plot model accuracy (with error bars) for a number of (new) random batches.

Vanishing Gradient

- Be creative and find a way to visualize the problem of vanishing gradients.
 - E.g. look at gradients for different sequence lengths